



Credit Card Default Prediction Using Machine Learning

CS 677 Final Project – Presentation | December 18, 2025 | Fnu Ashutosh (U01955320), Atharva Pande (U01985210), Kenji Okura (U01769019)

PROBLEM STATEMENT

Statement: Predict whether a credit card customer will default on their next payment, which enables proactive risk management and minimizing financial losses.

WHY? – THE BUSINESS IMPACT

Accurate predictions are needed to minimize financial loss for the credit card company. On average:

- Not predicting a default (false negative) costs \$10,000 per customer
- Incorrectly predicting a default (false positive) can cost \$200 per customer
- Class imbalance is 77.88 non-default to 22.12 default (3.52:1 ratio).

Because of the high cost of defaults, recall (or catching the defaults) a critical metric.

The model would need to prioritize catching true defaulters, while trying to minimize the chances of false positives. All to maximize profits.



Dataset

UC Irvine's Machine Learning Repository - Default of Credit Card Clients

QUICK STATS

- Sample Size: 30,000 customers
 - Training – test split: 80-20 (24,000 training, 6,000 test samples).
- Time period: April 2005 – September 2005 (6 months)
- Number of features: 23
 - All do not have null values (no missing data!)
 - All are integer data types
 - Examples include sex, limit_bal, education, payment status, etc.
- The target variable (if they defaulted next month) is binary (0 or 1).



UC Irvine
Machine Learning
Repository

WHY THIS DATASET?

- Real-world relevance with a sufficient size and many features. The default rate balance of 22.12 reflects real world distributions.

Exploratory Data Analysis (EDA)

The Dataset Structure, Basic Information, Statistics, and quick analysis

Shape and Size

30,000 samples
23 features
Totals to 690,000 data points

Data Types

The features are all integers
The target is a binary 0 or 1
Together they take 5.3 mb of space

Stats

We calculate the mean, standard deviation, min max etc. Some insights are:

- Credit Limit Range: NT\$10,000 - NT\$1,000,000
- Average Credit Limit: NT\$167,484
- Age Range: 21 - 79 years
- Average Age: 35.5 years

Missing Data

The source of data indicated there is no missing data. Analysis agreed with this.

Duplicate Data

There are 35 duplicate data.

Outliers

We tested outliers on select features (like age). We discover that LIMIT_BAL is most effected.

Exploratory Data Analysis (EDA)

Target Variable Distribution

Class Distribution


Why?

To apply correct evaluation metrics, decide on sampling strategies or setting realistic expectation


What we learned?

Moderately imbalanced, meaning class aware techniques is needed.

More than just testing for accuracy is needed.



Class Distribution - Percentage




Exploratory Data Analysis (EDA)

Analysis on Features

Demographic Features

Ex. Sex, education, marriage status, age.


For this case we calculated ratios of them for example.



Credit Limit Analysis

Ex. LIMIT_BAL

Find distribution, patterns or outliers




Payment Status

Ex. PAY_0....PAY_6

When analyzed with theme, we learn that payment behavior is improving



Payment Status

Ex. PAY_AMT1 to PAY_AMT6



Exploratory Data Analysis (EDA)

Correlation Analysis



Lessons Learned

We see high correlations for pay features, and demographics.

Once the grand correlation analysis we selected the highest correlations to compare the correlations.

See below.

Exploratory Data Analysis (EDA)

Key Findings

- The dataset is clean and ready for modeling .
- We do need to keep in mind we need to use techniques like class weightings and focus on Recall and F-1 score, not just accuracy
- The most important predictors are payment status and credit utilization, not demographics.
- The class imbalance and temporal nature of the data require special handling in modeling and feature engineering.

FEATURE ENGINEERING



Feature Engineering

Our Focus & Features

Our Focus

Based on our EDA insights, we engineered features focusing on:

- Credit utilization patterns (strong predictor of default)
- Payment behavior consistency (PAY features = top predictors)
- Temporal trends (bill/payment trajectories)
- Risk indicators (late payments, underpayment patterns)

High Level Features

1. Credit Utilization Features (3 features, i.e. average utilization across 6 months)
2. Payment Fraction Features (3 features)
3. Payment Behavior Features (2 features)
4. Temporal Trend Features (3 features)
5. Repayment Status Features (3 features)
6. Bill Stability Features (2 features)
7. Payment Consistency Features (2 features)
8. Risk Indicator Features (3 features)
9. Demographic Interaction Features (4 features, ex. Education-repayment)

We created 23 new features and verified they aren't missing values or infinite values.

EDA Engineered Features

Total engineered features: 25

| | count | mean | std | min | 25% | 50% | 75% | max |
|--------------------|------------|------------|------------|---------------|------------|------------|------------|--------------|
| util_mean | 30000.0000 | 0.3730 | 0.3519 | -0.2326 | 0.0300 | 0.2848 | 0.6879 | 5.3640 |
| util_max | 30000.0000 | 0.4950 | 0.4330 | -0.1000 | 0.0706 | 0.4307 | 0.9232 | 10.6884 |
| util_recent | 30000.0000 | 0.4238 | 0.4115 | -0.6199 | 0.0220 | 0.3140 | 0.8298 | 6.4550 |
| payfrac_mean | 30000.0000 | 185.5246 | 1664.3268 | 0.0000 | 0.0416 | 0.0986 | 1.5294 | 83050.6667 |
| payfrac_recent | 30000.0000 | 100.2243 | 2699.3173 | 0.0000 | 0.0352 | 0.0574 | 0.2942 | 298887.0000 |
| payfrac_min | 30000.0000 | 0.0461 | 0.1283 | 0.0000 | 0.0000 | 0.0009 | 0.0348 | 1.0276 |
| underpay_count | 30000.0000 | 3.4947 | 2.2027 | 0.0000 | 1.0000 | 4.0000 | 6.0000 | 6.0000 |
| overpay_count | 30000.0000 | 0.7960 | 1.1686 | 0.0000 | 0.0000 | 0.0000 | 1.0000 | 6.0000 |
| bill_trend | 30000.0000 | 4210.1761 | 33444.8078 | -1663964.0000 | -2348.0000 | 0.0000 | 5115.5000 | 504644.0000 |
| bill_oldtrend | 30000.0000 | 8141.3944 | 36159.8611 | -450018.0000 | -1829.0000 | 220.0000 | 11171.2500 | 1286872.0000 |
| pay_trend | 30000.0000 | 437.8990 | 20910.6391 | -845256.0000 | -436.0000 | 176.5000 | 1672.2500 | 503930.0000 |
| repay_max | 30000.0000 | 0.4387 | 1.3452 | -2.0000 | 0.0000 | 0.0000 | 2.0000 | 8.0000 |
| repay_mean | 30000.0000 | -0.1824 | 0.9822 | -2.0000 | -0.8333 | 0.0000 | 0.0000 | 6.0000 |
| months_late | 30000.0000 | 0.8342 | 1.5543 | 0.0000 | 0.0000 | 0.0000 | 1.0000 | 6.0000 |
| bill_std | 30000.0000 | 12077.7875 | 20302.1491 | 0.0000 | 1549.9352 | 4579.6622 | 14352.2596 | 647788.0511 |
| bill_range | 30000.0000 | 28850.2261 | 47606.9396 | 0.0000 | 4011.0000 | 11593.5000 | 33151.7500 | 1682177.0000 |
| zero_payment_count | 30000.0000 | 1.2299 | 1.7186 | 0.0000 | 0.0000 | 0.0000 | 2.0000 | 6.0000 |
| payment_cv | 30000.0000 | 0.9002 | 0.6533 | 0.0000 | 0.3798 | 0.7883 | 1.2982 | 2.4494 |
| high_util_flag | 30000.0000 | 0.2411 | 0.4277 | 0.0000 | 0.0000 | 0.0000 | 0.0000 | 1.0000 |
| deteriorating_flag | 30000.0000 | 0.4670 | 0.4989 | 0.0000 | 0.0000 | 0.0000 | 1.0000 | 1.0000 |
| chronic_late_flag | 30000.0000 | 0.0863 | 0.2809 | 0.0000 | 0.0000 | 0.0000 | 0.0000 | 1.0000 |
| age_util | 30000.0000 | 13.1022 | 13.3532 | -6.9588 | 1.0488 | 9.2619 | 22.1540 | 171.6493 |
| limit_per_age | 30000.0000 | 4829.5977 | 3712.5711 | 161.2903 | 1818.1818 | 4000.0000 | 6896.5517 | 26785.7143 |
| edu_repay | 30000.0000 | -0.2416 | 1.9525 | -12.0000 | -1.0000 | 0.0000 | 0.0000 | 16.5000 |
| gender_util | 30000.0000 | 0.5864 | 0.6066 | -0.4652 | 0.0483 | 0.4009 | 0.9338 | 10.7281 |

Missing values: 0

Infinite values: 0

All engineered features are clean (no missing/infinite values)

1. months_late : +0.3984 (↑ positive)
2. repay_max : +0.3310 (↑ positive)
3. chronic_late_flag : +0.3106 (↑ positive)
4. repay_mean : +0.2820 (↑ positive)
5. edu_repay : +0.2696 (↑ positive)
6. zero_payment_count : +0.1634 (↑ positive)
7. limit_per_age : -0.1599 (↓ negative)
8. overpay_count : -0.1263 (↓ negative)
9. util_mean : +0.1155 (↑ positive)
10. age_util : +0.1085 (↑ positive)
11. high_util_flag : +0.1080 (↑ positive)
12. underpay_count : +0.0984 (↑ positive)
13. gender_util : +0.0958 (↑ positive)
14. deteriorating_flag : -0.0903 (↓ negative)
15. util_recent : +0.0862 (↑ positive)

Feature Engineering Insights

General Insights


Top 3 engineered features by correlation with default:

1. months_late (count of months with late pay)
2. repay_max (max pay delay)
3. chronic_late_flag (>3 months late)

Observation (not just with top 3):

- Utilization and repayment features show strongest correlations
- Temporal trends capture payment behavior changes
- Risk flags effectively identify high-risk patterns

Data Viz Insights



- Higher values in risk features correlate with higher default rates
- Utilization features show clear risk gradients
- Payment behavior features distinguish defaulters from non-defaulters

Data preparation for model training

```
# Train-Test Split with Stratification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
57] ✓ 0.0s
```

- TRAIN-TEST SPLIT & PREPROCESSING

```
# Split with stratification to preserve class distribution (80/20 split per requirement)
X_train, X_test, y_train, y_test = train_test_split(
    X_engineered, y,
    test_size=0.2,
    random_state=RANDOM_STATE,
    stratify=y
)

print(f"\nDataset Split:")
print(f" Training set: {X_train.shape[0]} samples ({X_train.shape[0]/len(X_engineered)*100:.1f}%)")
print(f" Test set:     {X_test.shape[0]} samples ({X_test.shape[0]/len(X_engineered)*100:.1f}%)")
print(f" Features:    {X_train.shape[1]}")
58] ✓ 0.0s
```

```
Dataset Split:
Training set: 24,000 samples (80.0%)
Test set:      6,000 samples (20.0%)
Features:      48
```

```
Class Distribution:
Training default rate: 22.12%
Test default rate:      22.12%
Stratification preserved class balance
```

Features scaled (mean=0, std=1)

Sample scaled values (first 5 features, first sample):
Before: [1.6e+05 2.0e+00 2.0e+00 2.0e+00 3.3e+01]
After: [-0.05686623 0.80844039 0.18452304 0.85673912 -0.26455769]

MODEL DEVELOPMENT AND EVALUATION

XGBoost Baseline

General Insights

Model 1: Standard XGBoost Baseline - optimized for accuracy.

Some stats!

Accuracy: 0.8183

Precision: 0.6591

Recall: 0.3700

F1-Score: 0.4739

AUC (Area Under the ROC Curve): 0.7779

Missed Defaults (FN): 836

False Alarms (FP): 254

Caught Defaults (TP): 491

True Negatives (TN): 4419

```
print("\n1. Training Baseline XGBoost...")  
xgb_baseline = XGBClassifier(  
    n_estimators=100,  
    max_depth=6,  
    learning_rate=0.1,  
    random_state=RANDOM_STATE,  
    n_jobs=-1,  
    eval_metric='logloss'  
)
```

What is XGBoost?

XGBoost (aka eXtreme Gradient Boosting) uses gradient boosted decision trees (a learning algorithm that uses gradient descent) to reduce overfitting. XGBoost allows:

- Parallelism – speed up!
- Cache-awareness – speed up!
- Regularization (hyperparameter tuning)

XGBoost Baseline – Hyperparameter Tuning and impact

General Insights

Parameters tuning was done on: max_depth, learning_rate, n_estimators, and scale_pos_weight.

GridSearchCV and XGBClassifier were used together

Accuracy: 0.7747 (from 0.81)

Precision: 0.4923 (from 0.65)



Recall: 0.6051 (from 0.37)

F1-Score: 0.5429 (from 0.47)

AUC: 0.7819 (from 0.77)

Viz

- Max depth per tree: 3
- Each tree is learning from previous tree's errors per the graph
- Leaf values represent contribution to final prediction
- Feature splits chosen to minimize loss function




Most Important Features

We can analyze which features the decision trees use most for splitting decisions. For this model it's:

- repay_max, pay, etc.

Feature importance helps interpret what drives default prediction



XGBoost Cost-Effective

Model 2: Cost-effective XGBoost – false negatives are penalized more heavily. To enable cost-sensitive XGBoost, a “scale_pos_weight” was added with a 3.52:1 ratio considering the in-balance of classes.

Accuracy: 0.7620

Precision: 0.4715

Recall: 0.6285

F1-Score: 0.5388

AUC (Area Under the ROC Curve): 0.7757

Missed Defaults (FN): 493

False Alarms (FP): 935

Caught Defaults (TP): 834

True Negatives (TN): 3738

```
print("\n2. Training Cost-Sensitive XGBoost (scale_pos_weight=3.52)...")
imbalance_ratio = (y_train == 0).sum() / (y_train == 1).sum()
xgb_cost_sensitive = XGBClassifier(
    n_estimators=100,
    max_depth=6,
    learning_rate=0.1,
    scale_pos_weight=imbalance_ratio, # 3.52:1 ratio
    random_state=RANDOM_STATE,
    n_jobs=-1,
    eval_metric='logloss'
)
```

Random Forest Classifier

Model 3: Experiments with other models – Random Forest

Accuracy: 0.7758

Precision: 0.4943

Recall: 0.5923

F1-Score: 0.5389

AUC (Area Under the ROC Curve): 0.7802

Missed Defaults (FN): 541

False Alarms (FP): 804

Caught Defaults (TP): 786

True Negatives (TN): 3869

```
rf_model = RandomForestClassifier(  
    n_estimators=100,  
    max_depth=10,  
    min_samples_split=20,  
    min_samples_leaf=10,  
    class_weight='balanced', # Cost-sensitive learning  
    random_state=RANDOM_STATE,  
    n_jobs=-1  
)  
rf_model.fit(X_train_scaled, y_train)  
rf_time = time.time() - start_time  
  
# Predictions  
y_pred_rf = rf_model.predict(X_test_scaled)  
y_prob_rf = rf_model.predict_proba(X_test_scaled)[:, 1]  
  
# Evaluation  
rf_metrics = {  
    'model': 'Random Forest',  
    'accuracy': accuracy_score(y_test, y_pred_rf),  
    'precision': precision_score(y_test, y_pred_rf),  
    'recall': recall_score(y_test, y_pred_rf),  
    'f1': f1_score(y_test, y_pred_rf),  
    'auc': roc_auc_score(y_test, y_prob_rf),  
    'training_time': rf_time,  
    'cm': confusion_matrix(y_test, y_pred_rf)  
}
```

Random Forest Classifier – Hyperparameter tuning and impact

General Insights

Parameters tuning was done on: n_estimators, max_depth, and min_samples_split

GridSearchCV and RandomForestClassifier were used together

Accuracy: 0.7792 (from 0.7758)

Precision: 0.5096 (from 0.4943)

Recall: 0.5885 (from 0.6923)

F1-Score: 0.5410 (from 0.5389)

AUC: 0.7797 (from 0.7802)


Both models we did tuning on prioritize payment history (PAY_*) and credit utilization features

Viz

Shallow tree (less nodes) are simple but prevents over fitting

Deep tree is more complex but captures patterns

In this case shallow had better F1 score




Most Important Features

We can analyze which features the decision trees use most for splitting decisions. For this model it's:

- pay, repay_max etc.

Feature importance helps interpret what drives default prediction



Support Vector Machine

Model 4: More models we tested. This model is also cost sensitive

Accuracy: 0.7528

Precision: 0.4556

Recall: 0.6036

F1-Score: 0.5193

AUC (Area Under the ROC Curve): 0.7576

Missed Defaults (FN): 526

False Alarms (FP): 957

Caught Defaults (TP): 801

True Negatives (TN): 3716

```
svm_model = SVC(  
    C=1.0,  
    kernel='rbf',  
    gamma='scale',  
    class_weight='balanced', # Cost-sensitive learning  
    probability=True, # Enable probability estimates for AUC  
    random_state=RANDOM_STATE  
)  
svm_model.fit(X_train_scaled, y_train)  
svm_time = time.time() - start_time  
  
# Predictions  
y_pred_svm = svm_model.predict(X_test_scaled)  
y_prob_svm = svm_model.predict_proba(X_test_scaled)[:, 1]  
  
# Evaluation  
svm_metrics = {  
    'model': 'SVM (RBF Kernel)',  
    'accuracy': accuracy_score(y_test, y_pred_svm),  
    'precision': precision_score(y_test, y_pred_svm),  
    'recall': recall_score(y_test, y_pred_svm),  
    'f1': f1_score(y_test, y_pred_svm),  
    'auc': roc_auc_score(y_test, y_prob_svm),  
    'training_time': svm_time,  
    'cm': confusion_matrix(y_test, y_pred_svm)  
}
```

Kernel Functions Demo (on SVM)



We used Support Vector Machines to show comparing different kernel functions:

- Linear Kernel: For linearly separable data
- Radial Basis Function Kernel: For non-linear patterns
- Polynomial Kernel: For polynomial decision boundaries

Our Insights

We found that

- Linear: Slowest training
- RBF: Most flexible, fastest
- Polynomial: Medium speed, highest f-1 score



Logistic Regression (with SGD)

Model 5: We performed logistic regression with a SGD optimizer.

Accuracy: 0.7193

Precision: 0.4126

Recall: 0.6353

F1-Score: 0.5003

AUC (Area Under the ROC Curve): 0.7360

Missed Defaults (FN): 484

False Alarms (FP): 1200

Caught Defaults (TP): 843

True Negatives (TN): 3473

```
start_time = time.time()
sgd_model = SGDClassifier(
    loss='log_loss', # Logistic regression loss
    penalty='l2',
    alpha=0.0001,
    learning_rate='optimal',
    max_iter=1000,
    class_weight='balanced', # Cost-sensitive learning
    random_state=RANDOM_STATE,
    n_jobs=-1
)
sgd_model.fit(X_train_scaled, y_train)
sgd_time = time.time() - start_time

# Predictions
y_pred_sgd = sgd_model.predict(X_test_scaled)
# SGD decision_function for probability-like scores
y_prob_sgd = sgd_model.decision_function(X_test_scaled)

# Evaluation
sgd_metrics = {
    'model': 'Logistic Regression (SGD)',
    'accuracy': accuracy_score(y_test, y_pred_sgd),
    'precision': precision_score(y_test, y_pred_sgd),
    'recall': recall_score(y_test, y_pred_sgd),
    'f1': f1_score(y_test, y_pred_sgd),
    'auc': roc_auc_score(y_test, y_prob_sgd),
    'training_time': sgd_time,
    'cm': confusion_matrix(y_test, y_pred_sgd)
}
```

Gradient Boosting – We can skip

Model 6: The out of the box version of Gradient Boosting provided by sklearn

Accuracy: 0.7508

Precision: 0.4546

Recall: 0.6338

F1-Score: 0.5294

AUC (Area Under the ROC Curve): 0.7811

Missed Defaults (FN): 486

False Alarms (FP): 1009

Caught Defaults (TP): 841

True Negatives (TN): 3664

```
start_time = time.time()
gb_model = GradientBoostingClassifier(
    n_estimators=100,
    learning_rate=0.1,
    max_depth=3,
    min_samples_split=20,
    min_samples_leaf=10,
    subsample=0.8,
    random_state=RANDOM_STATE
)

# Gradient Boosting doesn't have class_weight, use sample_weight instead
sample_weights = np.where(y_train == 1, imbalance_ratio, 1.0)
gb_model.fit(X_train_scaled, y_train, sample_weight=sample_weights)
gb_time = time.time() - start_time

# Predictions
y_pred_gb = gb_model.predict(X_test_scaled)
y_prob_gb = gb_model.predict_proba(X_test_scaled)[:, 1]

# Evaluation
gb_metrics = {
    'model': 'Gradient Boosting',
    'accuracy': accuracy_score(y_test, y_pred_gb),
    'precision': precision_score(y_test, y_pred_gb),
    'recall': recall_score(y_test, y_pred_gb),
    'f1': f1_score(y_test, y_pred_gb),
    'auc': roc_auc_score(y_test, y_prob_gb),
    'training_time': gb_time,
    'cm': confusion_matrix(y_test, y_pred_gb)
}
```

Other Analysis

Learning Curves Analysis


We can use learning curves to detect underfitting and overfitting by showing how well the training and validation scores change with training size.

- Small train-validation gap shows good generalization (low overfitting)
- Converging curves show Model has learned effectively
- Large gap shows overfitting (model memorizes training data)
- Low scores for both show underfitting (model too simple)

Our Insights

Therefore, we can conclude that:

- XGBoost Cost-Sensitive it's overfitting
- Logistic Regression/SGD good fit



K-Fold Cross-Validation


This validation provides performance estimates across different training and validation data splits

- The 22.12% default rate was maintained
- Reduced variance in evaluation
- Low standard deviation shows stable, consistent performance
- High standard deviation shows model sensitive to data splits
- CV scores validate test set results (no lucky splits)

Our Insights

Therefore, we can conclude that:

- SVM is least stable and sensitive to data split
- Random forest still best
- Logistic Regression seems most stable



Final Findings



Grand Comparison

| Model | Accuracy | Precision | Recall | F1-Score | AUC | Missed Defaults (FN) | False Alarms (FP) | Caught Defaults (TP) | True Negatives (TN) |
|---------------------------|----------|-----------|--------|----------|-------|----------------------|-------------------|----------------------|---------------------|
| Random Forest | 0.7758 | 0.4943 | 0.592 | 0.5389 | 0.78 | 541 | 804 | 786 | 3869 |
| XGBoost Cost-Sensitive | 0.762 | 0.4715 | 0.629 | 0.5388 | 0.776 | 493 | 935 | 834 | 3738 |
| Gradient Boosting | 0.7508 | 0.4546 | 0.634 | 0.5294 | 0.781 | 486 | 1009 | 841 | 3664 |
| SVM | 0.7528 | 0.4556 | 0.604 | 0.5193 | 0.758 | 526 | 957 | 801 | 3716 |
| Logistic Regression (SGD) | 0.7193 | 0.4126 | 0.635 | 0.5003 | 0.736 | 484 | 1200 | 843 | 3473 |
| XGBoost Baseline | 0.8183 | 0.6591 | 0.37 | 0.4739 | 0.778 | 836 | 254 | 491 | 441 |

Findings

- Cost-sensitive models significantly improve recall (catch more defaults)
- Different models excel at different things
- XGBoost was most accurate, but had high false negatives
- Random forest does have higher F-1 score but has more false negatives

Focus Just On XGBoost

Cost Analysis

Baseline Model:

Missed defaults cost:

$$836 \times \$10,000 = \$8,360,000$$

False alarm cost:

$$254 \times \$200 = \$50,800$$

Total Cost: \$8,410,800

Cost-Sensitive Model:


Missed defaults cost:

$$493 \times \$10,000 = \$4,930,000$$

False alarm cost:

$$935 \times \$200 = \$187,000$$

Total Cost: **\$5,117,000**



More Analysis

From the graphic above we can see both models show similar discrimination ability (close AUCs)

And cost sensitive has much less false positives!

Choice!

XGBoosted Cost
Sensitive

Future

Limitations

1. Dataset Constraints
 - 6 month Taiwan-specific data may not generalize globally
 - Features: income, employment status, FICO scores are missing
2. Model Limitations
 - No deep learning methods explored
 - Threshold fixed at 0.5
3. Evaluation Scope
 - Test set performance only (no temporal validation)
 - Binary classification (no risk scoring tiers)
 - Cost assumptions simplified (\$10K/default, \$200/false alarm)

Future

1. Explore more datasets with:
 - International dataset with more features
 - More demographic features
2. Explore deep learning, ensemble methods (stacking, blending), neural networks (LSTM for temporal patterns), explainable AI (SHAP values for interpretability), time series feature extraction



THANK YOU!